

Examination 2020-10-16

Introduction to Parallel Programming (1DL530) – 5 hours

This is an examination **done at home**, so you can have your books and notes open when you take it. You also need to have access to a PC or server where you can compile and run C/C++ programs with Pthreads, OpenMP and MPI. Please refrain from searching the web for answers to these questions but, if you do, you **need to explicitly mention these sources**. Please avoid submitting hand-written text or low-quality pictures from mobile phones; instead **use an editor or word processor** for all your answers in text and, if possible, a drawing program for your answer to question 6. When you finish, you need to **upload your program for question 2 and a single PDF with your answers to the other questions** to your group in the Studium. Also, **remember to write your anonymous code in both your program and in the PDF with your answers**. Your submission **must be uploaded/received by 13:30**.

The exam contains **40** points in total and their distribution between sub-questions is clearly identifiable. To get a final grade of **5** you must score at least **34**. To get a grade of **4** you must score at least **28** and to pass the exam you must score at least **19**. Whenever in *real doubt* for what a particular question might mean, **state your assumptions clearly**.

DON'T PANIC!

1. Performance Metrics (2 + 1 + 1 + 4 = 8 pts total).

- Assume an application where the execution of floating-point instructions on a certain processor P consumes 60% of the total runtime. Moreover, let us assume that 25% of the floating-point time is spent in square root calculations. The design team of the next-generation processor P' believes that they can either improve the performance of all floating point instructions by a factor of 1.5 or alternatively speed up the square root operation by a factor of 15. From which design alternative would the aforementioned application benefit the most? Justify your answer by showing your calculations.
- (**Amdahl's Law**) Instead of waiting for the next processor generation, the developers of the application decide to parallelize its code.
 - (a) What speedup can be achieved on a 16-CPU system, if 90% of the entire program can be perfectly parallelized?
 - (b) What fraction of the code has to be parallelized to get a speedup of 10?
- (**Weak Scalability**) Algorithm A takes an integer k as input, uses $O(k^2)$ workspace, and has cubic (i.e., $O(k^3)$) time complexity. Let $T_p(k)$ be the execution time of A to solve a problem of size k with p processes. It is known that $T_1(4000) = 12$ minutes. Assuming perfect weak scalability, what is the expected execution time for $k = 64000$?

2. **Parallel Programming Using Pthreads (4 + 2 + 2 = 8 pts total).** Conway's Game of Life should be known to you from Assignment 3 of this (and last) year's course. Please refer to the text of that assignment if you do not recall the details of the game. That assignment also came with a sequential implementation of the game given in file `Game_of_Life.c`, which took as arguments the size of the array and the number of generations (steps) in the game.
- (a) Parallelize the part of the code between the two `gettimeofday()` calls using Pthreads. You will need to submit your code (the resulting `Game_of_Life_par.c` file), together with a brief description of what its parallelization involved and the main changes you did. Please refrain from doing unnecessary changes in the remaining functions of the file.
 - (b) Did you need to use locks or any other synchronization mechanisms and why?
 - (c) Using a departmental server (or your laptop/PC if it is powerful enough to do this task), measure and report the speedup your program achieves when running with 2, 4 and 8 threads on a 1024×1024 array using 4000 steps.
-

3. **Parallel Programming (1 + 1 + 2 + 2 = 6 pts total).** Consider the code below:

```
int    locks[100];
float  data[1000000];

#pragma omp parallel for
for (i = 0; i < 1000000; i++) {
    lock(locks[g(i)%100]);
    f(data[g(i)]);
    unlock(locks[g(i)%100]);
}
```

Assume that the function `f` only accesses its argument, may read and write it, and does not do any synchronization internally. Give brief answers to the following questions:

- (a) How many threads can be active at any time?
 - (b) Assuming locking is free, what is the maximum speedup that this program can achieve?
 - (c) Is the program data race free or not? Explain your answer.
 - (d) Assuming locking is free, is it possible that the loop executes sequentially because of locking?
-

4. **Parallelization Using OpenMP (6 pts total).** The code below takes an array and "rotates" it by shifting all its elements to the left and appending the first element.

```
void rotate(int n, double *a) {
    for (int i = 0; i < n - 1; i++) {
        double tmp = a[i];
        a[i] = a[i+1];
        a[i+1] = tmp;
    }
}
```

Parallelize the code of this function using OpenMP and submit your solution. Make sure you properly handle the dependencies that exist in this code.

5. **OpenMP Programming (6 pts total).** A programmer has parallelized a `foo` function using OpenMP pragmas as follows:

```
int foo(int n, int *a, int *b, int *c, int *d) {
    double tmp, total = 0;
    #pragma omp parallel
    {
        #pragma omp sections
        {
            #pragma omp section
            {
                #pragma omp for
                for (int i = 0; i < n; i++) {
                    a[b[i]] += b[i];
                    total += b[i];
                }
            }
            #pragma omp section
            {
                #pragma omp for
                for (int i = 0; i < n; i++) {
                    tmp = c[i];
                    c[i] = d[i];
                    d[i] = tmp;
                    total += c[i];
                }
            }
            #pragma omp for
            for (int i = 0; i < n; i++) {
                total += d[i];
            }
        }
    }
    return total;
}
```

If you think there are any issues (i.e., errors, omissions in the pragmas, or pragmas which are missing) with this code, explain why they are issues and suggest a (good) way to fix them.

6. **MPI Programming (3 + 3 = 6 pts total).** In lecture 11, we examined *MPI Collective Communication* and saw different ways that communication can be depicted using trees. We also discussed the functions `MPI_Scatter` and `MPI_Gather`. Suppose `comm_sz` = 8 and `n` = 16.

- Draw a diagram that shows how `MPI_Scatter` can be implemented using tree-structured communication with `comm_sz` processes when process 0 needs to distribute an array containing `n` elements.
- Draw a diagram that shows how `MPI_Gather` can be implemented using tree-structured communication when an `n`-element array that has been distributed among `comm_sz` processes needs to be gathered onto process 0.

Good luck !