

# Exam for Real Time Systems

2024 Jan 11

Wang Yi

## Important Instructions:

1. No course material or computer/calculator are allowed, only a pen and a dictionary.
2. Please mark which course you are registered for:

☐ 5hp (1DT063)

You need to solve problems 1–4 only  
(100p)

☐ 10hp (1DT004)

You need to solve *all* problems  
(150p)

3. For each problem, a number of choices/statements are given, that may be correct or wrong.

You are asked to mark only the **correct** ones.

If you mark all correct choices for each problem, it will give you all points of the problem; otherwise each wrong choice (marked) will result in "negative points" proportionally.

.....  
**Problem 1.** 1. (5p) *Why is it so difficult to estimate the Worst-Case Execution Time of a real-time task?*

- ( ) *it may have too many loops*
  - ( ) *it may need too much memory in execution*
  - ( ) *it may have too many instructions*
  - ( ) *it may consume too much energy*
  - ( ) *the platform may support multi-tasking*
  - ( ) *the platform may have caches*
  - ( ) *it is impossible to decide whether a program terminates or not*
  - ( ) *the platform may have a multicore processor*
2. (5p) *A real-time programming language like Ada compiler should support:*
- ( ) *object orientation*
  - ( ) *multi-tasking*
  - ( ) *implementation of real-time tasks*
  - ( ) *task synchronization*
  - ( ) *user-defined scheduling policies*
3. (5p) *Resource servers can be useful for*
- ( ) *upgrading the hardware platform of an existing system*
  - ( ) *isolating applications of different types*
  - ( ) *improving the average response times of soft real-time tasks*
  - ( ) *guaranteeing the deadline of event-driven real-time tasks*



4. (5p) When RMS is used, a task set is non-schedulable if its utilization is
  - ☐ over 100%
  - ☐ over 69.3%, but less than 100%
  - ☐ less than  $n * (2^{1/n} - 1)$  where  $n$  is the number of tasks
5. (5p) The advantages of fixed-priority scheduling include:
  - ☐ it is stable
  - ☐ it can be used to implement predictable systems
  - ☐ it is the best policy for optimizing resource utilization
  - ☐ it is the best policy for achieving the schedulability of real-time tasks
  - ☐ it is the best policy for improving the average response time of real-time tasks
  - ☐ it can be used to implement safety-critical systems
6. (5p) Assume a set of sporadic sets  $(C_i, D_i, T_i)$  where deadline  $D_i \leq T_i$ . If DMS is used, the task set is schedulable if
  - ☐  $(\sum C_i/D_i) \leq 100\%$
  - ☐  $(\sum C_i/T_i) \leq 100\%$
  - ☐  $(\sum C_i/D_i) \leq n * (2^{1/n} - 1)$  where  $n$  is the number of tasks
  - ☐  $(\sum C_i/T_i) \leq n * (2^{1/n} - 1)$  where  $n$  is the number of tasks
7. (5p) A system with periodic tasks may contain dependent tasks with input/output relations, which causes release jitters. Assume that the dependent tasks have the same period and task B depending on task A is defined as follows: the computed result of a job of task A is used as the input of a job of task B (and thus task B suffers a release jitter due to the varying response times of task A).  
Is it possible to avoid this and thus avoid release jitters?
  - ☐ this is not possible at all
  - ☐ yes, deliver the result of a job when it is completed
  - ☐ yes, deliver the result of a job at the worst-case response time estimated
  - ☐ yes, deliver the result of a job by the end of each period
  - ☐ yes, deliver the result of a job by the given deadline if the tasks are schedulable
8. (5p) For a set of periodic task scheduled on a single processor, what are the differences between DMS and EDF?
  - ☐ EDF is stable; DMS is non stable
  - ☐ EDF is non stable; DMS is stable
  - ☐ DMS uses fixed priorities; EDF uses dynamic priorities
  - ☐ EDF is good for resource utilization; DMS is not as good as EDF
  - ☐ EDF is optimal; DMS is not
  - ☐ EDF is optimal; DMS is also optimal
9. (5p) Assume that a set of periodic tasks with deadlines equal to periods is schedulable on a single processor system with fixed-priority scheduling according to a user-defined fixed-priority assignment. Then the task set is also schedulable using
  - ☐ non-preemptive RMS
  - ☐ RMS
  - ☐ non-Preemptive EDF
  - ☐ EDF
  - ☐ non-preemptive DMS
  - ☐ DMS



- ☐ *FIFO*
- ☐ *Shortest Job First Scheduling*

10. (5p) When you use CAN, you should

- ☐ *make sure that messages sent from different nodes have different identities*
- ☐ *know that sending messages with the same identities by different nodes may lead to errors*
- ☐ *send messages with periodic tasks*
- ☐ *assign the priorities of messages according to the periods of the sending tasks*
- ☐ *make sure that the priority ordering of messages is a total order according to their id's with the lower id number, the higher priority for messages*
- ☐ *send messages with receiver's address*

**Problem 2.** (20p) Assume a CAN bus running at 1 Mbits per second, connecting 3 stations (nodes) A, B and C. There are 12 messages of fixed size (200 bits each), with 12 identities (1 to 12), sent by the nodes as follows:

1. On node A, there are four tasks each sending messages with identities 1,2,3,4 at most every 1ms, 2ms, 2ms and 7ms respectively.
2. On node B, there are four tasks each sending messages with identities 5,6,7,8 at most every 8ms, 10ms, 18ms and 20ms respectively.
3. On node C, there are four tasks each sending messages with identities 9, 10, 11, 12 at most every 30ms, 31ms, 43ms and 50ms respectively.

The worst case transmission delay (i.e. the time period from queuing to completed message transmission) for message with identity 1 is bounded by:

- ☐ 0.10ms
- ☐ 0.30ms
- ☐ 0.50ms
- ☐ 0.70ms
- ☐ 0.90ms
- ☐ 1.10ms

The worst case transmission delay for message with identity 6 is bounded by:

- ☐ 0.2ms
- ☐ 0.40ms
- ☐ 0.60ms
- ☐ 0.80ms
- ☐ 1.00ms
- ☐ 1.20ms
- ☐ 2.20ms
- ☐ 3.20ms
- ☐ 4.20ms
- ☐ 5.20ms





**Problem 3.** (15p) The two standard operations  $P$  and  $V$  on semaphores are implemented according to the following pseudo-code:

P(S)	V(S)
(1) Disable-interrupt; (2) if S.counter > 0 (3) then S.counter -- 1; (4) else { (5) insert(current-task, S.queue); (6) schedule(); (7) Enable-interrupt	(8) Disable-interrupt; (9) If non-empty(S.queue) (10) then { (11) next-to-run := get-first(S.queue); (12) insert(next-to-run, Ready-queue); (13) schedule(); (14) else S.counter ++ 1; (15) Enable-interrupt

To implement HLP (Immediate Priority Inheritance Protocol), the following changes may be made:

- ( ) between line (3) and (4), it should be extended by "save the priority of the current-task and raise it to the ceiling of S"
- ( ) between line (4) and (5), it should be extended by "save the priority of the current-task and raise it to the ceiling of S"
- ( ) between line (14) and (15), it should be extended by "restore the priority of the current-task to the previously saved priority"

**Problem 4.** (15p) Assume an event-driven task with computing times  $C_i$  and deadlines  $D_i$  (assume  $C_i \leq D_i$ ). The jobs of the task may arrive sporadically but the minimal distance between two jobs is  $T_i > D_i$ . Design a polling server for the task with period  $T$  and capacity  $C = C_i$  such that no deadline is violated when RMS is used for scheduling. Assume that  $C \leq T$  and  $R$  is the response time estimated using RMS analysis when the server is considered as a periodic task. Further, assume  $R \leq T$ . Identify correct statements:

- ( )  $T \leq D_i$  is necessary
- ( )  $T \leq D_i$  is sufficient
- ( )  $T + T \leq D_i$  is necessary
- ( )  $T + T \leq D_i$  is sufficient
- ( )  $T + R \leq D_i$  is necessary
- ( )  $T + R \leq D_i$  is sufficient
- ( )  $T + C \leq D_i$  is necessary
- ( )  $T + C \leq D_i$  is sufficient





**Problem 5.** (20p) Assume a set of sporadic tasks running on a platform of  $N$  processors each with the same speed  $S$  (or processing capacity). Assume that the task set is schedulable with global scheduling. Identify the correct statements:

- ☐ The task set with some tasks removed is also schedulable.
- ☐ The task set with the computing time of some tasks reduced is also schedulable.
- ☐ The task set is also schedulable on the same platform with more memory.
- ☐ The task set is also schedulable on the same platform with more caches.
- ☐ The task set is also schedulable with a platform of more processors of speed  $S$ .
- ☐ The task set is also schedulable with a platform of  $10 * N$  processors of speed  $S/10$ .
- ☐ The task set is also schedulable with a single powerful processor of speed  $N * S$ .
- ☐ The task set is surely feasible with partitioned scheduling.
- ☐ The task set is surely feasible with partitioned scheduling (allowing job-splitting) if its utilization is below  $N * 69.3\%$ .

**Problem 6.** (10p) Assume a set of periodic tasks with periods as deadlines and utilizations as follows: 0.80, 0.90, 0.9, 0.4, 0.7, 0.7, 0.3, 0.8, 0.9, 0.2, 0.1, 0.2, 0.1

If you partition the tasks onto a multiprocessor platform and then run the respective scheduling policy, what is the least number of processors you need?

- ☐ 9-12 with EDF
- ☐ 5-8 with EDF
- ☐ 1-4 with EDF
- ☐ 9-12 with RMS
- ☐ 5-8 with RMS
- ☐ 1-4 with RMS



**Problem 7. (20p)**

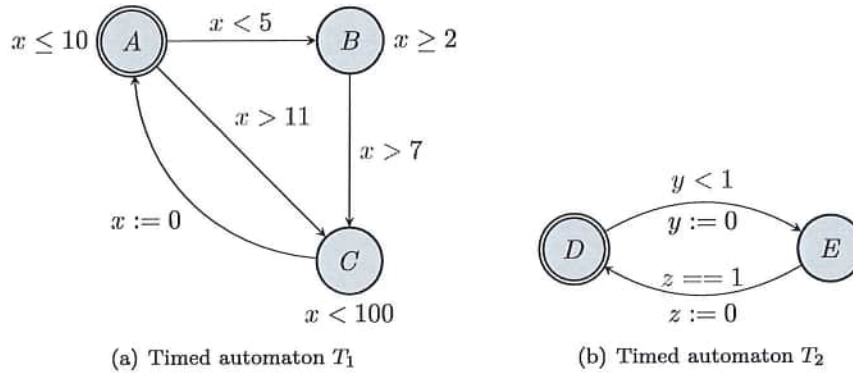


Figure 1: Two timed automata in which  $x$ ,  $y$  and  $z$  are clocks. Note that both automata are separate systems, they do *not* run together.

Take a look at the timed automaton  $T_1$  in Figure 1(a). Identify correct statements:

- ☐ Location  $B$  is reachable
- ☐  $x$  can be below 5 at location  $C$
- ☐ There are no edges that can be removed without changing the set of reachable states
- ☐ It satisfies  $E \leftrightarrow T_1.B$
- ☐ It satisfies  $E \leftrightarrow (T_1.C \text{ and } x > 5)$
- ☐ It satisfies  $E \leftrightarrow (T_1.C \text{ and } x < 5)$
- ☐ It satisfies  $A \leftrightarrow x \geq 1$

Take a look at the timed automaton  $T_2$  in Figure 1(b). Identify correct statements:

- ☐ It is possible that  $T_2$  deadlocks
- ☐ No matter what,  $T_2$  will eventually deadlock

